

# EV Battery Health Monitor: Industry Standards and Best Practices

## 1. Industry Standards for EV Battery Telemetry Data

**Data Formats & Protocols:** Unlike some IT domains, EV battery telemetry lacks a single universal data format. Each automaker often defines proprietary CAN bus messages or API fields for battery data. However, there are emerging guidelines. For example, China's **GB/T 32960** standard mandates that EVs transmit battery telemetry (e.g. pack voltage, current, temperature, State of Charge) to a national cloud platform in real time <sup>1</sup>. This is used for safety monitoring and warranty analysis, and it essentially forces a standardized data schema at least within China. In Europe and the US, standards are looser: vehicles typically log data via the CAN bus and OBD-II protocols (with extended PIDs for EV-specific metrics), or via automaker-specific telematics APIs. An industry trend is toward open data access APIs – for instance, Tesla provides a **Fleet Telemetry API** that streams vehicle data at **500 ms intervals** (2 Hz) <sup>2</sup>. This is quite granular for cloud telemetry and indicates the level of detail modern EVs can report. Other OEMs (BMW, Ford, etc.) expose data through connected vehicle services (e.g. BMW's CarData or Ford's FordPass APIs), though often at lower frequencies or on-demand requests (e.g. a REST call for the latest status every few minutes). In summary, **there is no single schema used by all manufacturers** – but key fields and concepts are common across the industry.

**Key Metrics Tracked:** EV battery telemetry focuses on a core set of health and performance metrics that have become *de facto* standards:

- **State of Charge (SoC):** The battery's charge level (% full). This is the most critical metric for both drivers and fleet managers, essentially the EV's "fuel gauge" <sup>3</sup>. SoC is used to plan trips and charging – vehicles often send SoC updates periodically or when changes exceed a threshold.
- **Battery Energy (kWh):** The absolute energy remaining in the pack, in kilowatt-hours. This complements SoC by showing actual energy available. Fleet software uses this to match vehicles to routes (ensuring a vehicle has enough kWh for a trip) <sup>4</sup>.
- **Voltage & Current:** Pack voltage (overall) and current (flow in or out of battery). These are fundamental electrical measurements. For example, **pack current** indicates charge/discharge rate (positive during charging, negative during driving). Sudden drops in voltage under load can indicate battery stress or faults. Many systems derive **power** (kW) from voltage\*current.
- **Temperatures:** Most EVs report battery temperature (often the average or several sensor readings). High battery temperatures can indicate thermal issues; cold batteries may have reduced performance. Some advanced systems log the maximum and minimum cell temperatures for more granular monitoring. (Tesla's BMS data, for instance, includes multiple thermistor readings and flags like `BatteryHeaterOn` which indicates active thermal management <sup>5</sup> <sup>6</sup>.)
- **State of Health (SoH):** A measure of battery degradation, expressed as a percentage of original capacity. SoH is essentially the ratio of current max energy capacity to the initial capacity <sup>7</sup>. Industry convention considers ~80% SoH as end-of-life for automotive use <sup>8</sup>. Automakers usually calculate SoH in the BMS and may expose it in service data or via apps (e.g. Nissan and GM show

“battery health” bars). MOBI (an industry consortium) recently released a standard defining SoH formally to encourage consistent reporting <sup>9</sup> .

- **Cycle Count & Charge History:** Some systems log how many charge cycles the battery has undergone or the cumulative energy throughput. These help in warranty and lifespan analysis. (For example, Tesla vehicles internally track total kWh charged/discharged; not always exposed to end-users, but used in diagnostics.)
- **Remaining Range (Estimated):** Given SoC and recent consumption, cars report estimated range. This is more of a computed metric, but it's often part of telemetry (e.g. `est_battery_range` in Tesla API gives projected miles remaining <sup>10</sup> ). It's crucial for fleet operations to prevent vehicles from stranding <sup>11</sup> .
- **Location & Usage Context:** Though not battery metrics per se, telemetry often pairs battery data with GPS location, speed, and usage patterns. Location is important for fleet charging logistics <sup>12</sup> and also for environment context (e.g. uphill driving might correlate with high current draw). Many OEM telematics systems (including Tesla's) include GPS and even environmental data like ambient temperature.
- **Alerts and Fault Codes:** Batteries have diagnostic trouble codes and warning flags (over-temperature, cell imbalances, etc.). Fleet platforms consider these “vehicle alerts” critical <sup>13</sup> . For instance, an alert that battery voltage is dropping abnormally fast could signal a failing cell. Standards like OBD-II and ISO 15118 (for charging) define some of these codes, but much is proprietary. A good battery monitoring app could listen for such flags to label “events” (e.g. an over-temperature event).

**Sampling Rates & Data Volume:** In passenger EVs, the battery management system (BMS) monitors cells in real time (often at 10 Hz or faster internally), but not all that data is uploaded. Telematics units commonly send aggregated data at 1 Hz or lower, possibly event-triggered. As noted, Tesla streams at 2 Hz to its cloud API <sup>2</sup> , which is quite high frequency. Other OEMs might send data every few seconds or on state change (to save bandwidth). **Standard logging for many fleet/IoT applications is around 1 Hz (once per second) per parameter**, which balances granularity with data size. At 1 Hz, a single car can generate ~86,400 data points per day for one metric. If an EV logs dozens of metrics, that's on the order of millions of data points per vehicle-year. Real-world fleet data confirms this scale: an AWS automotive telemetry blog mentions **“millions of vehicles, each generating hundreds of metrics every second,”** illustrating the massive scale when data from many cars aggregates <sup>14</sup> .

Because of this volume, **data retention policies** are important. In industry, raw high-frequency data might be kept for a shorter period (weeks or months) and then down-sampled or summarized for long-term storage. For example, a backend might retain 1 Hz data for 3 months, then keep only 1-minute averages for a year, and perhaps daily statistics thereafter. Regulatory requirements can influence retention: in some jurisdictions (e.g. under Chinese GB/T rules), manufacturers must store certain battery data for the life of the vehicle. Elsewhere, warranty and safety analysis often drive retention (e.g. keep data for at least the warranty period). Since EV batteries are expensive components, manufacturers tend to err on the side of keeping data that might explain a future failure. **Data privacy laws** (like GDPR in Europe) also affect telemetry retention – personal data (including location tied to a vehicle/VIN) should not be kept longer than necessary without consent <sup>15</sup> . A portfolio project can mention it plans to implement configurable retention (perhaps using the database's retention policy or partition drop features) to mirror these practices.

**Industry Protocols and APIs:** Modern EVs typically send telemetry via cellular networks using either MQTT-like publish/subscribe (common in IoT) or HTTPS to cloud endpoints. For instance, Tesla vehicles maintain a persistent secure connection for the streaming API. Many OEMs integrate with cloud IoT platforms: **MQTT** is

popular for its lightweight overhead. Some industry initiatives aim at standardization: the **MQTT Sparkplug** spec or **Open Telematics** define how to structure IoT messages (key-value pairs with timestamps, etc.), which could be adopted for EV data. Additionally, the **Open Charge Point Protocol (OCPP)** is a standard for EV charger-to-cloud communication – not for in-car data, but relevant if your app ever ingests charging station telemetry. To align with best practices, you might structure your data as timestamped records with vehicle ID, metric name, value, and units. This is essentially a **time-series schema**. A common approach is the **sensor data model**: e.g. a table with columns `[timestamp, vehicle_id, metric_name, value]` or, more efficiently, one table per metric category. In fact, TimescaleDB's docs and many IoT systems favor **wide tables per device** (one row per timestamp per device, with columns for each metric) for performance <sup>16</sup> <sup>17</sup>. Major players likely use binary proprietary formats on the wire (for efficiency), but convert to standardized schemas in their cloud data stores.

**Takeaway:** For your project, you can confidently include fields like `timestamp`, `vehicle_id`, `SoC`, `pack_voltage`, `pack_current`, `temperature`, etc., knowing these align with industry norms. Ensure the units are clear (e.g. °C for temperature, V for voltage, A for current, % for SoC). Including **State of Health** in your schema (even if you simulate it) will demonstrate awareness of battery lifecycle monitoring. Also, consider implementing basic support for standard file formats (if doing batch upload): e.g. CSV or JSON logs with time-series data. There's no single "EV telemetry file format" standard publicly used, but using CSV with headers like `time, soc_percent, voltage_V, current_A, temp_C` would be very familiar to employers and can be mapped to how data comes from CAN loggers or fleet CSV exports.

## 2. Tech Stack Validation (TimescaleDB vs Alternatives in 2024–25)

Your chosen stack of **React/Next.js (frontend)**, **FastAPI (Python backend)**, and **PostgreSQL+TimescaleDB (database)** is very much defensible for a 2024/2025 project – both for practicality and in the job market. Let's examine the database choice in particular, as that's crucial for time-series data:

**TimescaleDB on PostgreSQL:** TimescaleDB (an extension to PostgreSQL) has matured into a popular time-series solution. It brings automatic time-based partitioning, compression, and fast aggregation functions to the rock-solid Postgres engine. This means you get the reliability and flexibility of a relational database (ACID transactions, standard SQL, rich indexing) while handling time-series efficiently. For an IoT-like use case (EV telemetry), Timescale is absolutely realistic. In industry, we see many companies choosing Postgres/Timescale for time-series when they want to avoid maintaining a separate specialized DB. TimescaleDB's performance is strong for moderate-to-high volumes: e.g. internal benchmarks have shown **1,000x faster queries and ~90% storage reduction** versus vanilla Postgres for time-series workloads <sup>18</sup>. With compression enabled on older data, TimescaleDB achieved "over 90%" size reduction in real deployments, and one case study reported a **97% reduction in storage** for IoT sensor data (from 64 GB/day uncompressed down to ~2 GB/day) <sup>19</sup>. This compression not only saves space but *improves* analytic query speed by using columnar storage for historical chunks. In terms of ingestion rate, Timescale can handle a high throughput on decent hardware. It's been measured ingesting on the order of **millions of records per second** in benchmarks <sup>20</sup> – far beyond what your single-node MVP will need. In short, TimescaleDB can easily handle a fleet's data if scaled properly (and your MVP likely will be orders of magnitude smaller, which is fine).

**Job Market & Industry Use:** Knowing TimescaleDB is a plus for many data engineering roles, even if not explicitly required. It demonstrates you understand time-series data modeling and Postgres internals. While

some companies use alternatives (InfluxDB, DynamoDB for IoT, etc.), the concepts overlap. In 2024/25, using SQL for time-series is actually *trendy* again, due to tools like Timescale, ClickHouse, and DuckDB that blend analytical capabilities with familiar query languages. There are also managed services: for example, Azure offers Timescale as an option in their PostgreSQL service, and Crunchy Data has a cloud Timescale option – indicating demand. Thus, you can confidently say “I chose TimescaleDB because it allows high-performance time-series processing on PostgreSQL, which is a stack many companies are comfortable with.”

**Comparison – InfluxDB:** InfluxDB was the go-to open-source time series DB of the last decade. It’s still widely used for metrics and IoT. InfluxDB (v1.x) had a custom SQL-like query (InfluxQL) and v2 uses Flux script language; it’s purpose-built for time-series with easy retention policies and built-in downsampling. Influx can be very fast for writes and compression is good. However, it introduces an entire new technology to learn/maintain, whereas Timescale rides on Postgres (leveraging your existing SQL skills and any ORM, etc.). Performance-wise, the two are comparable for many workloads – some benchmarks show Timescale outperforming Influx for complex queries especially <sup>21</sup>. For example, one test found Timescale **188% faster** than Influx on certain heavy analytical queries <sup>21</sup>. Influx might edge out Timescale for raw ingestion in a single node scenario (it’s optimized for append-only writes), but Timescale catches up when data volume grows and queries become complex (thanks to indexed relational storage and parallel query plans). For a portfolio project, Timescale is a great choice because it avoids the operational overhead of another DB and still proves you can handle time-series data expertly. In a job interview, you can mention InfluxDB as an alternative you evaluated, which shows awareness. If asked “why not Influx?”, you can note that Timescale’s use of SQL allows easier integration with other application logic and that for moderate scale (e.g. a few million points), Timescale performance is more than sufficient (indeed, **Timescale outperformed some cloud time-series services like AWS Timestream by orders of magnitude in tests** <sup>17</sup>).

**Comparison – Prometheus:** Prometheus is a time-series database, but aimed at infrastructure monitoring (think server CPU metrics). It’s pull-based and focuses on recent data with efficient in-memory storage and a custom query language (PromQL). Prometheus excels at monitoring thousands of metrics with retention of maybe weeks – but it’s not meant as a long-term historical store for IoT data. Data can be “remote written” to long-term storage, but usually not to itself. For an EV telemetry system, Prometheus would not be an ideal primary DB for user data because of its ephemeral nature and lack of relational querying. It *could* be used if you were building a real-time monitoring dashboard only (like visualizing latest metrics with alerts), but since you want data labeling and analysis, a traditional DB (SQL or NoSQL) is more appropriate. However, knowledge of Prometheus is valued in DevOps roles, and you can mention that your time-series skills would translate if monitoring pipelines come up.

**Comparison – Cloud TS Services:** Cloud providers have services like **AWS Timestream**, **Azure Time Series Insights (TSI)**, or using **Azure Data Explorer (Kusto)** for time-series, and Google tends to use BigQuery or Bigtable for similar needs. These managed services are great for real production at scale: for example, AWS Timestream is serverless and handles retention tiers automatically. In a professional project, one might choose Timestream or Azure TSI to avoid managing the DB. But for your local portfolio (and budget constraint), using TimescaleDB in Docker is perfectly valid. In fact, by using an open-source DB, you show more hands-on skill (cloud services sometimes abstract away the details). If interviewers ask about scaling, you can discuss how you would consider moving to a cloud service or clustering Timescale if the data grew into the billions of rows. You can also mention that Timescale has a proven track record: a benchmark by Timescale themselves showed **6,000× higher insert throughput and 150× cheaper cost** compared to

Amazon Timestream for a write-heavy workload <sup>17</sup> . Even if those numbers are vendor-biased, it underlines that Timescale on decent hardware is very powerful.

**Comparison – Other Alternatives (ClickHouse, Druid, etc.):** In the time-series/big-data space, **ClickHouse** has emerged as a star for analytics. It's a blazing fast columnar OLAP database that handles time-series and event data at scale (originally by Yandex, used by Cloudflare, etc.). ClickHouse could absolutely handle EV telemetry (it's known to ingest millions of rows per second and execute sub-second aggregations on billions of records). Some companies use it for telemetry dashboards and off-line analysis. However, it has a different SQL dialect and is slightly more complex to operate (distributed cluster setup, etc.). For a single-node project, ClickHouse would be overkill and less convenient than Timescale. **Apache Druid** and **Apache Pinot** are other specialized stores, often used in realtime analytics dashboards (e.g. Uber's analytics, LinkedIn feeds). They shine for events and time-indexed records with aggregation, but similarly are complex to set up for a small project and usually require a cluster and Zookeeper, etc. **QuestDB** and **TDengine** are newer open-source time-series DBs claiming top performance, but they are niche and not widely asked for in jobs (showing familiarity is nice but not necessary).

Overall, **TimescaleDB is a well-balanced choice**: it's industry-grade enough to impress (and it's used in domains like industrial IoT, energy management, smart cities, etc.), yet simple enough to manage on a laptop. Its use of SQL aligns with the *full-stack developer* angle of your project – you can write complex queries, joins (e.g. joining event labels with raw data), and even use it for app authentication or metadata if needed (since it's still Postgres underneath). This one-stack approach (Python + Postgres) is realistic in many smaller companies or prototyping teams.

**Backend and Frontend choices:** FastAPI and React/Next.js are also very contemporary. FastAPI (with Uvicorn/ASGI) is known for being fast and easy to write, and it's Python – great for integrating data science (Pandas/NumPy) for any analysis tasks. Many data engineering teams use Python for APIs, so that skill is valuable. On the frontend, React with Next.js shows you know how to build modern interactive web apps. Plotly.js for charts is a good choice for time-series visualization; it's widely used in scientific and financial dashboards. Employers in automotive might not care specifically which charting library you use, but they **do** care that you can present data in a clear, interactive way on the web. Using TypeScript on the frontend is also a plus (shows you adhere to best practices in web development).

**Running on WSL2:** Developing on Windows via WSL2 (Ubuntu 22.04) is common now, and it demonstrates you're comfortable with a Linux environment. The stack you chose (Docker, Postgres, etc.) runs on Linux containers by default, so WSL2 is giving you a native-like Linux kernel to work with. In 2024, Docker Desktop with WSL2 integration is stable and performant. There are virtually no incompatibilities in using TimescaleDB on WSL2 versus a real Linux server – it's the same code running. The main consideration is performance tuning: by default, WSL2 will consume as much memory as needed up to a limit (usually 50% of host RAM). If you load a lot of data, monitor the `vmmem` usage. You can cap it via a `.wslconfig` if needed (e.g. to prevent your system from swapping). Another WSL2 tip is to ensure your Postgres data is stored in the Linux filesystem (e.g. under `/var/lib/postgresql` in the container). Avoid mounting a Windows directory for the database files, as that would introduce slow file I/O due to cross-OS overhead <sup>22</sup> <sup>23</sup> . Keeping everything within WSL2's filesystem yields near-native performance – WSL2 file access is fast from Linux side, and Docker will do this by default unless you override volumes to point to `/mnt/c` or similar. In short, **this stack works very well in WSL2**, and many developers at companies use WSL2 similarly to mimic a Linux server environment on a Windows dev machine. You likely won't hit any unusual bugs. Just develop as if you're on Ubuntu (VS Code Remote-WSL is excellent for this, as you're using).

**Database in Docker vs Native:** One more note – you’ll be running PostgreSQL/Timescale inside Docker. That’s fine for development and demo. TimescaleDB doesn’t require special kernel tuning beyond standard Postgres configs (which the Docker image usually sets reasonably). For production, a company might run the DB on a dedicated server or a managed service, but using Docker in WSL2 for now is perfectly acceptable. If you find the need to persist data between runs, Docker volumes or bind mounts can be used (again, target the WSL2 filesystem for those). In interviews, you can mention that developing via Docker containers in WSL2 allowed you to ensure the app is portable and works cross-platform – e.g. it could run on a Linux VM in the cloud just as easily.

**TimescaleDB vs others – conclusion:** Yes, your PostgreSQL/TimescaleDB choice is absolutely defensible in 2024/2025. It aligns with industry trends of leveraging scalable SQL systems for time-series (often termed “HTAP” – hybrid transactional/analytical processing). It will also allow you to demonstrate real SQL prowess (writing a complex query to e.g. compute average degradation or annotate anomalies with SQL window functions will impress more than showing a simple NoSQL usage). Just be prepared to discuss alternatives (Influx, etc.) to show you did your homework – but it sounds like you already intend to, given the question prompts.

### 3. Architecture Best Practices for Telemetry (Real-Time vs Batch, Ingestion, Modeling)

Designing an EV telemetry system requires balancing **real-time data processing** (for immediate alerts/monitoring) and **batch or historical processing** (for analytics, training ML models, and so on). Companies in automotive and IoT often adopt a **Lambda architecture** (or similar) – a speed layer and a batch layer:

- **Real-Time Ingestion & Processing:** This is typically handled via event streams. Modern stacks use message brokers like **Apache Kafka** (or its cloud equivalents like AWS Kinesis, Azure Event Hubs) to ingest high-frequency data from vehicles. Each vehicle’s telematics unit publishes messages (JSON or binary) to a topic (e.g. “vehicle\_123/battery”). A streaming pipeline then processes those in real time – for example, filtering or triggering alerts (high temperature alert, etc.). The AWS blog on EV analytics shows an architecture where vehicles stream data through AWS IoT Core into an **MSK (Managed Kafka)** cluster <sup>24</sup>. From Kafka, data was fed into an analytics engine (in that case, OpenSearch for real-time dashboards and alerts) and simultaneously to S3 for long-term storage <sup>25</sup>.
- **WebSockets/MQTT vs REST:** The industry leans toward push-based streaming for live data. A WebSocket or MQTT subscription can deliver data to front-end dashboards with low latency. REST APIs, on the other hand, are used for on-demand queries of historical data or aggregate stats (e.g. “give me the last 7 days of battery temps”). For your MVP, it’s sensible to start with batch file ingestion or a simple REST upload (simulating a vehicle log being uploaded). This covers the core functionality without needing a full live pipeline. You can then outline a “Phase 2” where a streaming component (like a Kafka or Redis pub/sub, or even just FastAPI sending data to clients via WebSocket) would be added to show real-time monitoring. Simply being able to discuss how you’d incorporate streaming will demonstrate architecture acumen. For instance, you could say: *“In the future, we could insert a Kafka pipeline to handle hundreds of vehicles streaming concurrently – the FastAPI server would produce events to Kafka, and a consumer would write to TimescaleDB, while another consumer could feed a live WebSocket dashboard.”* Even if you don’t implement it fully, understanding that pattern is key.

- **Batch Processing & Analytics:** This involves processing large chunks of data (e.g. nightly jobs computing daily summary metrics, or retraining an ML model on accumulated data). In automotive contexts, batch processing might be used for things like computing long-term battery degradation trends or running heavy analytics that aren't urgent. Many companies dump raw telemetry into a data lake (Hadoop or cloud storage like S3/ADLS) and then run Spark jobs or use Databricks to crunch it. In the AWS architecture reference, the S3 storage was for offline business intelligence queries and ML training <sup>25</sup>. For your project, batch processing is essentially your main mode (since you'll likely simulate or upload log files). Optimize this path to be as fast as possible so it *feels* interactive. Techniques: use vectorized libraries (NumPy/Pandas) or COPY into the database instead of row-by-row inserts if ingesting a large file. TimescaleDB can ingest data quickly, but for very large CSVs consider using PostgreSQL's COPY command for efficiency.

**Data Modeling & Partitioning:** Best practices for time-series modeling include partitioning data by time and perhaps by another dimension like vehicle\_id. TimescaleDB's **hypertables** automatically partition by time (and you can specify a secondary "space" partition, e.g. vehicle\_id, if you want). Partitioning by time makes it efficient to drop or compress old chunks and improves insert performance by localizing writes. With Timescale, you'll define a hypertable on the main telemetry table. For example, a table `telemetry(vehicle_id, ts, soc, voltage, current, temp, ... , PRIMARY KEY(vehicle_id, ts))` turned into a hypertable partitioned by `ts` (time) in say 1-week chunks would be a solid design. You might also create a smaller table for events/anomalies (each labeled event with a time range, vehicle\_id, and event type). This table will be much lower volume and can be easily joined with the telemetry data when querying by time ranges.

**Indexing:** It's typical to index time-series data on the time column (which Timescale does for you) and on identifiers like vehicle\_id. This allows querying one vehicle's data fast. In a multi-tenant fleet system, an index on (vehicle\_id, ts) or partitioning by vehicle would prevent one vehicle's data from affecting another's query. Since your use case is portfolio scale (maybe a handful of vehicles), a single hypertable partitioned by time is fine. If you wanted to showcase extra knowledge, you could partition by time *and* by vehicle in Timescale, which effectively creates sub-partitions per vehicle – but that's more relevant at huge scale (hundreds of vehicles, multi-billion rows).

**Ingestion at Scale:** Companies dealing with thousands of vehicles often use a pipeline like: *Vehicle -> IoT Hub (or MQTT broker) -> Stream processing -> Database*. For example, **Geotab** (a telematics provider) built a custom ingestion platform using Kafka Connect to handle billions of records per day <sup>26</sup>. Another pattern is using **Apache NiFi or Flink** for ingesting and processing streaming data (applying filters or transformations before storage). Since you aim for local and cost-free, you won't spin up Kafka on AWS, but you can mention a lightweight alternative: *"For a simple demo, I use Python async tasks to simulate streaming; in a real system I'd use Kafka or Redis Streams to decouple ingestion from processing."* Kafka is indeed a buzzword in data engineering roles, so demonstrating familiarity helps. You might even include a diagram in your documentation with Kafka in dotted lines as a future addition.

**Real-Time vs Batch in UI:** Real-time data implies perhaps showing the latest data updating live (like a moving chart). Batch/historical analysis implies being able to scroll/zoom through past data. The best practice is to separate concerns: a **streaming service** (could be part of your FastAPI using WebSocket) handles pushing new points to subscribed clients, whereas the **REST API** handles queries like "give me data from time A to B" or "give me all labeled anomalies". This way, historical queries can be optimized (possibly

hitting pre-aggregated data). If you implement just the batch part now, you can still outline the real-time extension.

**WSL2-Specific Docker Considerations:** Running Docker in WSL2 is generally smooth, but there are a few *gotchas/optimizations* worth noting (and mentioning to show you encountered and solved them): - *File I/O*: As noted earlier, avoid mounting Windows paths for database storage. Instead, use Docker volumes or WSL2 paths. Accessing Windows files from Linux is **much** slower due to filesystem translation overhead <sup>22</sup>. For example, if you mistakenly had your Postgres data directory on `C:\...` mounted into the container, you'd see poor write performance. Best practice is to keep data in the Linux FS and, if needed, expose results to Windows by copying out or using the `\\wsl$` share for occasional access. - *Memory and CPU*: By default, Docker Desktop with WSL2 will let the VM (and thus containers) use resources on demand. This is great for performance (you get essentially the whole machine), but can sometimes cause Windows to page if you use too much. You can configure `.wslconfig` to set memory limit if you want to ensure, say, Docker doesn't use more than 8GB and choke Windows UI. Also, note that shutting down Docker doesn't automatically free WSL2 memory (it will shrink over time or when WSL is restarted). If you demonstrate on a laptop with limited RAM, you might mention you tuned these settings to ensure a smooth demo. - *Networking*: Docker on WSL2 avoids the old Hyper-V NAT issues; you can usually access services at `localhost:port` directly. Just be aware if you expose ports (like FastAPI on 8000, or Postgres on 5432) those are accessible from Windows host as well. For security in a demo, it's fine. - *Developer workflow*: Using VS Code Remote WSL means all your file operations are on Linux side, which is good. One thing to highlight: you can build the Docker images from WSL seamlessly. Some users get tripped up by building from Windows vs WSL – but you're doing it right by staying in WSL. Essentially, you are mimicking how many cloud deployment environments work (Linux container, Linux OS). That's a positive to emphasize: **"I developed and tested everything in a Linux environment (via WSL2) to ensure compatibility with real deployment targets."** Employers appreciate that over someone who only knows Windows-specific tooling.

**Time-Series Schema Best Practices:** Another aspect is how to design the schema for flexibility. Some systems use a **wide table** (one column per metric), others use a **key-value model** (one row per metric per timestamp). For a manageable set of known metrics (SoC, voltage, etc.), a wide table is simplest and fastest. You'll have columns like `soc_percent FLOAT, voltage FLOAT, current FLOAT, temp FLOAT, ...`. Timescale can handle hundreds of columns, but you likely have on the order of 5-20 important ones – which is fine. If you foresee wanting to add metrics dynamically or handle different vehicle models with different sensors, a more normalized model could be used (e.g. a separate `battery_measurements` table with columns `metric_type, value`, and a join to a timestamp table). That's likely overkill here. Mentioning that you considered both models (and chose one) shows design thinking. The wide-table approach aligns with typical BMS log CSVs, and allows simple SQL queries (e.g. `SELECT time, voltage, current FROM telemetry WHERE vehicle_id=X` yields a time-series easily).

**Continuous Aggregations & Downsampling:** TimescaleDB offers continuous aggregates, where it can pre-compute summaries (e.g. 1-minute averages of each metric) in the background. In production, this is often used to speed up queries of long time spans (instead of reading billions of raw points). For MVP, you may not need it, but you could configure one as a demonstration (say, a continuous aggregate for daily min/max/avg of key metrics). This is a talking point in interviews: *"The database is set to continuously compute daily statistics per vehicle, which is a pattern used in industry to enable quick dashboard loads and to implement retention (raw data beyond 6 months can be dropped if we have daily summaries)."* Even if you don't actually drop data, it shows you know how to design for scale.



**Example Architecture Diagram:** To visualize a best-practice architecture, consider this reference from AWS's EV telemetry solution. It shows vehicles streaming data into a Kafka pipeline, real-time indexing for dashboards, and a data lake for batch analytics. In a similar vein, your design could have a streaming component feeding the **time-series database (TimescaleDB)** for live queries, while also archiving data to cheap storage for offline analysis. For now, your "stream" might just be simulated by batch uploads, but the principles are identical.

The embedded diagram (from AWS) illustrates how a fully managed pipeline might look: IoT devices (EVs) → IoT Core (ingestion) → Kafka (stream buffer) → OpenSearch (for real-time search/visualization) and S3 (for historical data storage). In your case, **TimescaleDB will play the role of both the real-time and historical store**, since it can handle recent inserts and also store months of data (especially with compression). This simplifies the architecture for an MVP. It's worth noting to interviewers that *for the scale of a demo or a small deployment, a single database can suffice, but the design could evolve to a multi-tier pipeline as in that AWS reference for very large fleets*. Demonstrating that you're aware of cloud-scale patterns but chose an appropriate level of complexity for the project is key.

## 4. Employment Market Analysis for Automotive/IoT Data Roles (2024–25)

To ensure your project aligns with what employers are looking for, let's analyze the in-demand technologies and skills in the automotive/EV IoT space:

**Popular Tech Stacks & Tools:** Companies in the EV and IoT domain often use a mix of cloud platforms and data tools. Common elements include:

- **Cloud Infrastructure:** AWS and Azure are heavily used in automotive backends. For example, AWS IoT Core + Kinesis or Azure IoT Hub + Stream Analytics show up in job descriptions. If you can relate your project to these (even conceptually), that's a plus. For instance, "If deployed on AWS, my design could use AWS IoT for ingestion and Timescale on AWS (via RDS or Aurora) for storage" – this shows cloud fluency.
- **Streaming Frameworks:** As discussed, Kafka is a big one. Even if you don't implement it, being able to talk about producing/consuming telemetry events, partitioning topics by vehicle, and using consumer groups for scale will resonate. Some job listings mention "*experience with event streaming (Kafka, Pulsar or similar)*". Knowing Kafka's role in decoupling producers and consumers is good. If you have time, you could set up a small Kafka in Docker just to log a few events and mention it, but it's optional.
- **Databases:** Time-series databases (InfluxDB, Timescale, DynamoDB, Cassandra) or general NoSQL (MongoDB) sometimes appear. By using Timescale, you effectively cover SQL + time-series. TimescaleDB experience specifically might not be asked unless the company uses it, but it falls under "PostgreSQL experience" which is very widely valued. If a job uses Influx or Cassandra, they know a strong Postgres/Timescale engineer can learn those. Emphasize that you have designed schema and optimized queries for time-series – that skill is transferable. Also, mention any familiarity with big-data warehouses if you have (e.g. "I'm aware how this data could later be stored in Spark or Snowflake for big picture analysis"). Not that you need to use those in the project, but awareness is good.
- **Data Processing & ML:** For more data science-oriented roles (or companies like Tesla that do a lot of in-house analytics), skills like **PySpark/Spark, Pandas**, or even **MATLAB/Simulink** for battery modeling can come up. In your project, you are focusing on data engineering and labeling rather than model training, which is fine. However, showcasing a bit of data analysis (maybe a Jupyter notebook that computes a degradation curve or an anomaly detection example using your data) could impress. It says: not only can I collect data, I can derive insights. If time permits, you could train a simple anomaly detection (even a z-score outlier or ARIMA) on your time series

and mention it. But since the prompt suggests not building full ML models, this could be left as a discussion point in interviews (a roadmap item).

- **Programming Languages:** You're using Python and TypeScript – both highly relevant. Python is dominant in data roles; TypeScript/JS is common for any full-stack or front-end aspect. In automotive, embedded roles use C/C++, but for cloud and data, Python/Java/Scala are common. You might get asked if you know any Java or C++ because some automotive data systems integrate with embedded code; if you have exposure, mention it. If not, focusing on Python is fine.
- **DevOps & Containers:** Docker experience is a big plus as most companies containerize their applications. Your use of Docker Compose to tie together the DB, API, and UI shows you have DevOps awareness. Employers might value that you can set up a local environment quickly. Also, using Git for version control (make sure your portfolio code is on GitHub with a clear README) is essential. If you automate some deployment (like using GitHub Actions to build/test your app), that's cherry on top – it demonstrates CI/CD knowledge.
- **TimescaleDB vs other TSDBs in Job Market:** While Timescale isn't a household name like MySQL or Mongo, it has a niche following. Some job postings might explicitly mention it, especially in IoT or monitoring companies. More commonly they might say "experience with time-series data or TSDBs (InfluxDB, Timescale, etc.)". If you come across a posting from an industrial IoT firm, knowing Timescale could actually match their stack (Timescale has case studies in energy, manufacturing, smart cities <sup>16</sup>). But even if not, just advertise it as PostgreSQL experience with specialization in time-series – that covers both generic and specific skill.

**Skills to Emphasize for Battery/Automotive Tech Positions:**

- **Domain Knowledge:** This is huge. If you can talk about batteries – e.g. how high C-rate charging affects longevity, what thermal runaway is, how state of health is estimated – you will stand out against generic data engineers. It shows genuine interest in the field. So definitely emphasize what you've learned: mention terms like "coulomb counting for SoC estimation" or "cell balancing" or "battery degradation curves". You don't have to go deep into electrochemistry, but knowing the basics (like that extreme SOC or temperatures accelerate degradation) will let you have an intelligent conversation with an automotive team. Your project inherently helps here, since you're labeling thermal events or charge anomalies, etc., which implies you know those are important things to watch.
- **IoT and Embedded Integration:** Even though your work is cloud-side, understanding how the data is generated at the vehicle is valuable. For example, mention that you're aware of how a BMS collects cell measurements and how that might be packaged into CAN messages. If you've played with any hardware (like reading OBD-II data via a dongle), bring that up. Some automotive employers love hobby projects like a Raspberry Pi CAN bus sniffer or using an OBD-II adapter to log data from a car. It shows initiative and practical skills.
- **Full Stack Ability:** Since you want to show full-stack competence, highlight that you built everything from the database to the API to the UI. This is attractive to smaller companies or teams that need versatile engineers. You can say *"I'm comfortable switching between writing SQL queries for data analysis and designing a React frontend for visualization"*. Many roles, even if primarily data engineering, appreciate someone who can at least create internal dashboards or tooling to interface with the data. Your project being end-to-end demonstrates that nicely.
- **Analytics and Visualization:** In automotive analytics roles, being able to visualize data for engineers or decision-makers is important. Emphasize your Plotly dashboard as a feature, not just an afterthought. Perhaps mention that you followed good UX principles for interactive plotting (zooming, filtering) because effective visualization of time-series helps in diagnosing vehicle issues quickly. Also, mention if you styled it in an "automotive-like" manner (maybe using dark themes similar to dashboard displays, etc.) – it subtly shows attention to UI/UX, which is a bonus skill.
- **Testing and Reliability:** Automotive industry puts a premium on safety and reliability. It might be beyond scope to implement, but you could mention designing the system with reliability in mind – e.g. "using transactions in the database to ensure label data integrity" or "considering how to handle missing data or out-of-order data (since networks can drop or delay packets)". If you note these concerns and perhaps write

a bit about them in documentation, it shows an engineer's mindset for robust systems. - **Certifications/Frameworks:** There aren't specific "battery software" certifications, but there are relevant ones: **AWS Certified Solutions Architect** (if you highlight cloud integration knowledge), or **Certified Kubernetes Administrator** (if dealing with deployments). Those are general though. In automotive, there are standards like **ISO 26262 (functional safety)** – not directly relevant to your project unless you were building on-vehicle code, but awareness might impress if it comes up. Another angle: **SAE standards** for telematics – e.g. *SAE J1939* is for heavy vehicle CAN, *SAE J1979* covers OBD-II PIDs, *ISO 15118* covers EV charging communication. Dropping an acronym or two like "we could parse data akin to OBD-II PID for battery SoC" shows extra mile. As for formal certs, if you had any (say, an AWS cert or a course in battery management), mention them, but if not, it's okay. Experience and projects weigh more. - **Value of TimescaleDB Experience:** Timescale is still niche, but any PostgreSQL-heavy role will value your deep knowledge of it. You can spin it this way: by using Timescale, you had to learn about database performance tuning (since you likely encountered how to use indexes, how to partition, etc.). That is broadly useful. If an employer uses say DynamoDB for IoT, you can relate concepts like partition keys and time-indexing from your project. Timescale also gave you exposure to SQL for analytics (with `time_bucket` and `window` functions perhaps), which is directly applicable to many data analyst/engineer tasks. So yes, mention it as a positive experience: *"I worked extensively with a time-series extension on Postgres, which taught me a lot about efficient data modeling and query optimization for time-indexed data."* - **Cross-Platform & WSL2:** Many automotive companies develop on Linux (for embedded) or have cloud systems on Linux servers, but also have developers using Windows (especially for CAN analysis tools which often run on Windows). Showing that you bridged these by using WSL2 is great. It basically says *"I use Windows as an interface but I'm developing in Linux under the hood."* That means you can fit into teams regardless of dev environment. Definitely, in interviews, mention that you're comfortable with Linux command line, shell scripting, etc., thanks to WSL2. Some interviewers might not know WSL2 in detail, so you can frame it simply: *"I did all my development in a Linux environment (Ubuntu) on my machine to mirror production, using Docker to containerize the app."* That covers it. Cross-platform skill also means if someone asked you to run something on a Mac or directly on Ubuntu server, you could.

**Geographic Considerations:** You noted US primarily, with some EU remote roles. In the US, companies like Tesla, Rivian, Lucid, the Big 3 (GM, Ford) all are hiring for EV data roles. They often list skills like Python, SQL, cloud, Hadoop, etc. European companies (BMW, VW, etc.) similarly look for big data and cloud experience. Remote roles might include telematics startups or fleet analytics companies. By covering Python, SQL, web dev, and IoT knowledge, you're hitting the checkboxes for a lot of these positions. If you wanted to further tailor, you could incorporate minor things like localization (just mention that you considered units conversion or timezone handling – since global fleets deal with timezones and metric/imperial units). But that's a minor detail.

**Soft Skills & Showcasing:** Don't forget to present this project in a way that highlights problem-solving and communication. Perhaps prepare a short "pitch" about it: the problem it solves (monitoring and labeling battery health events), why that matters (preventing fires, optimizing range, etc.), and how you built it. Employers love when candidates are passionate and articulate about their projects. Given the depth you're adding, you'll be able to handle deep-dive questions on anything from "How do you detect an anomaly?" to "Why did you choose WebSockets over polling for real-time updates?" etc. Use this to your advantage to steer conversation towards things you know well.

## 5. Implementation Guidance and Project Features

Now, focusing on actually building this MVP in 4 weeks – here are some practical guidelines and features to make it successful and aligned with industry expectations:

**Realistic Data Simulation:** Simulating data that *looks and feels* real is important for demonstrating your system. EV battery telemetry has recognizable patterns: - **Discharge during Driving:** SoC will decrease roughly linearly (with some curve) during a trip. Current will fluctuate (spikes when accelerating, negative spikes (regeneration) when braking). Temperature might rise slowly as the battery is used. Voltage will drop under load (voltage generally correlates with SoC for a lithium-ion battery – high SoC ~ high voltage, but under heavy current draw voltage dips). - **Charge Cycles:** When charging, SoC goes up (usually quickly at first if fast charging, then taper off near 100%). Current is positive (into battery) and might start high and then ramp down as you reach high SoC (in DC fast charging, a **CC-CV profile** means constant current then constant voltage). You can simulate a fast charge session by having current start at, say, +200 A (if simulating a supercharger) and gradually drop to +50 A as SoC goes from 20% to 80%. Temperature often increases during fast charge because of internal resistance heating. - **Idle:** When car is parked, SoC stays flat (or very slowly drains), current ~0, voltage might slowly settle. Temperature will tend toward ambient. These baseline periods are good to include so that your anomaly detection isn't always "on the move" – real data has a lot of boring idle time too. - **Anomalies to Include:** Think of a few story-driven events: - *Thermal event:* e.g. simulate a cooling system failure where battery temp suddenly rises faster than normal during a high-power event. This could be an anomaly label "overheating". You can implement a rule: if battery temp > X °C and rising > Y °C/min, flag it. - *Capacity fade:* over weeks of data, slightly reduce the effective full capacity so that SoH goes from 100% down to say 95%. This could be shown by the difference between the energy that corresponds to 100% SoC at the start vs end of dataset. It's subtle, but you can maybe calculate an estimated SoH and show a trendline. Label one point as "capacity dropped 5% over 6 months" in a summary. - *Sensor glitch:* maybe a single timestamp where voltage or current jumps unrealistically (due to a sensor error). Label it as "telemetry glitch" – shows you know data isn't perfect. - *Rapid degradation event:* if you want a dramatic anomaly, simulate that at one point the battery lost a chunk of capacity (could be from cell failure). This might look like SoC suddenly dropping faster than usual or the voltage curve shifting. It's rare in real life, but sometimes a bad cell can cause sudden range loss. You could label "cell failure suspected". - *Charging anomaly:* e.g. the vehicle didn't charge when expected (perhaps current stayed zero even though it was plugged in) – could indicate a charger fault. Or the charging was unusually slow. If you have charger data (like pilot signal, etc., which you likely won't, but you can simulate as simple as "expected to charge to 80% but only reached 50%"). - **Data Volume:** Decide how much data to generate. Even a week of 1Hz data is ~604,800 seconds. That might be fine to load if your system is efficient, but perhaps start with a few days per vehicle. You could simulate, say, 10 vehicles for 10 days each at 1Hz – that's ~8.6 million rows, which is a lot but Timescale can handle that on a decent PC, especially if you batch insert. However, for a demo, you don't need millions of points to show trends. A clever approach is to simulate at maybe 5 or 10 second intervals for long stretches, and higher frequency only for specific interesting periods. This reduces total points but still gives a sense of resolution when zoomed in. Since it's a portfolio, you want interaction to be smooth. A plot of 100k points can still be interacted with, but a plot of 5 million might choke the browser. Timescale can aggregate on the backend, but Plotly in the browser should not get megabytes of raw data at once. **Solution:** page or aggregate data for plotting. You could implement an API like `/data?start=...&end=...&resolution=1s or 1min` that returns downsampled points if the range is large. This mimics how professional dashboards work (coarser data for long ranges, raw data for zoomed-in ranges). - **Open Datasets & APIs:** While you likely simulate data, it's great to reference open data. For example, NREL or NASA have battery lifecycle datasets (those are more about charge/discharge cycles of

single cells in lab, not whole EVs, but you could use them for validating your SoH calculations). Geotab published a study analyzing data from 6,300 EVs and found an average 2.3% annual degradation <sup>27</sup> – you can cite such stats in your project README to justify why you simulate X% degradation. There's a **Kaggle dataset of EV charging** (likely the one you found <sup>28</sup>) which contains fields like SoC and Voltage during charging sessions. Perhaps you can use a snippet of that to compare your simulation. If that dataset is small, you could even include it as a sample in your system (e.g. have an "Import sample data" button). - Another source: Some EV enthusiasts share data logs on forums (e.g. Tesla API data). If you find any real log (even a short one) you could run it through your system to show it works with real data too. But that's optional.

**Visualization Standards in Automotive Dashboards:** While you're not building an in-car display, you can take inspiration from automotive UI: - **Clear Gauges/Indicators:** It might be nice to have a visual battery icon or gauge for current SoC when you inspect a specific time. For example, when the user clicks on a point, show a battery icon filled to that SoC%. Little touches like that make it memorable. - **Time-Series Graphs:** Typically, engineers looking at telemetry use line charts with multiple axes or multiple overlaid lines. Plotly can do multiple y-axes (e.g. plot SoC and current on one graph, or temperature on a second axis). Make sure to label axes with units (volts, amps, etc.). Use color coding consistent with automotive norms (maybe blue for SoC, red for temperature (as red often denotes heat/danger), green for current when charging, etc.). - **Interactive Features:** The ability to zoom and pan is standard (Plotly has it built-in). A useful addition is a **brush or range selector** – e.g. a mini-map below the chart where you can select a window. Plotly's range slider could be enabled for the x-axis to allow quick zooming by dragging it. - **Annotation/Labeling UI:** Since one of your goals is anomaly labeling, design a way to mark regions. A common UX is click-and-drag to highlight a time range, then a dialog pops up to tag that range (e.g. "Label as: Overheating event"). Alternatively, you could allow clicking a start point and end point for the event. Once labeled, the region could be shaded or an icon placed. This interaction will *really* stand out in an interview because it shows you built a custom, data-intensive UI element. It demonstrates front-end skill, understanding of user needs (e.g. an engineer wants to mark the exact period of an anomaly), and integration (since labels go to the backend DB). - **Automotive Aesthetics:** Perhaps use a dark theme with bright colored lines, as many dashboards (and indeed engineering tools) use dark backgrounds for plotting (less eye strain, and car interiors often have dark-themed displays). If using Next.js, you could pick a design library or simple CSS to make it look professional – spacing, fonts, etc. Even though it's a tech demo, polish matters. It subconsciously tells employers "I can deliver user-friendly tools, not just raw tech." - **Responsiveness and Performance:** Ensure the UI can handle, say, displaying a day's worth of data smoothly. If you need to, aggregate points on the server (Timescale can do `time_bucket('1 minute', ts)` to downsample). This way, if a user looks at a week overview, you show 72460 = 10080 points instead of ~604k points (1Hz). Plotly also has decimation features for large datasets which you might explore. - **Example Visuals:** If you can find references, consider how Tesla's own service tool might display battery data or how third-party tools like **TeslaMate** or **Grafana dashboards for EVs** do it. Some fleet platforms show a timeline with events annotated (e.g. "charging session here, idle here, driving here"). You could mimic that by coloring the background of the chart or using markers for events (like a plug icon when charging). This is beyond basic requirements, but even one or two such touches will show domain knowledge (e.g. you know when a car is charging vs driving based on current flow, and you reflect that in the visualization).

**Open Data Integration:** To add realism, you might incorporate external data like: - Weather API: Battery performance is affected by ambient temperature. If you tag each data point with outside temp (either simulated or fetched from a weather API for the location & time), you could then have an insight like "this

range drop occurred because it was -5°C outside.” This might be too much to implement fully, but you could simply simulate different ambient temps and note the effect. Or integrate a free API (if any) to get temperature by city/time – careful with API keys/cost though. - Maps: Probably out of scope, but if you had location data, plotting routes or using a map could be neat. That said, focusing on battery metrics is primary. - Known EV datasets: As mentioned, Kaggle or academic datasets of EV fleets. Even referencing them in your documentation (like “the simulation is informed by statistics from XYZ paper or dataset”) adds credibility. For example, **Geotab’s 2020 study**: they found **on average 2.3% annual degradation** and that climate and high DC fast charge usage can accelerate it <sup>27</sup>. You could say you simulate a vehicle that fast-charges a lot to show a worst-case degradation scenario.

**WSL2 + Docker Workflow Best Practices:** While developing, some tips (which you may know, but also to mention in documentation for others or in interview): - Use **docker-compose** YAML to define all services (DB, backend, frontend). This makes it one-command to start the whole stack. It’s impressive to say “just run `docker-compose up` and the whole app comes up networked together”. - For a smoother dev experience, you might not containerize the frontend during active development (you can do `npm run dev` on Windows or WSL side for hot-reload). But ensure you can produce a production build and serve it, perhaps via Next.js’s integrated server or a simple Nginx. In Docker, you might end up with 3 containers: `frontend` (Next.js server or static bundle), `backend` (FastAPI), `db` (Postgres+Timescale). Configure the networking so that FastAPI can talk to DB (usually compose does this with a default network). - **Volume mounts for code:** In dev mode, you can mount your FastAPI app directory into the container for rapid iteration (so you don’t have to rebuild the image on every code change). VS Code can even attach to the container. This is a bit advanced, so not required, but shows DevOps skill if done. - Document any needed environment variables in a `.env` (like database URL, etc.), and show that you’re not hard-coding secrets. Security is a small concern since it’s local, but good habits count. - Windows-specific: sometimes file permission issues arise with volumes on WSL2 (due to different UID/GIDs). If you encounter that, one fix is to add `:delegated` or proper `chown` in Dockerfile. Mentioning you solved such an issue (if it came up) indicates practical troubleshooting ability.

**Features to Stand Out to Employers:** Ultimately, to stand out, your project should **tell a story** and solve a problem that an automotive company cares about. Here are features and use-cases that showcase deep domain knowledge: - **Anomaly Detection & Labeling:** As planned, this is your core. Make it shine by including several types of anomalies (thermal, performance, charging, etc.) and providing a clean interface to label them. Perhaps implement some basic auto-detection (rule-based) to pre-label obvious ones and allow the user to adjust the labels. Being able to talk about how you’d integrate an ML model for anomaly detection in the future (e.g. “we could train a model using these labeled events to detect anomalies automatically on the edge device”) will show foresight. - **Degradation Analysis (SoH over time):** Include a view or report that tracks State of Health. For example, compute the effective battery capacity each month from the data (maybe by looking at how much energy it takes to charge from X% to Y%, etc.). Plotting **SoH vs time** or **range degradation vs time** would resonate strongly – it’s a key concern for EV owners and manufacturers. If you can simulate one vehicle that’s treated kindly (slow charges, moderate climate) vs one that is abused (lots of fast charges, extreme temps) and show their SoH diverging, that would be an insightful addition. It demonstrates understanding of what factors influence battery life. - **Remaining Useful Life (RUL) projection:** You said you’d mention it in roadmap – that’s wise. If you have SoH trend, you can extrapolate roughly to say “if trend continues, battery will reach 80% SoH (end-of-life) in 5 years”. This is RUL in essence. In interviews, you can discuss how one might build a predictive model for RUL using historical data and maybe external features. This aligns with current industry interest – e.g. startups and OEMs are very interested in second-life and end-of-life predictions to manage battery warranties and

recycling. - **Fleet perspective:** If you simulate multiple vehicles, you can show a fleet dashboard. For instance, a table of all vehicles with their current SoC, SOH, maybe last anomaly, etc. Fleet managers want at-a-glance status. Even if your focus is one vehicle at a time in graphs, having an overview page listing vehicles ranked by health or in need of attention (e.g. "Vehicle 7 has high degradation or had 3 anomalies this month") would impress. This is something Ampcontrol or other fleet software do – turning raw data into actionable info. - **Integration with Charging Infrastructure:** This might be beyond scope, but you could mention it. E.g. "Our system could integrate with charger data (via OCPP or charging session logs) to correlate battery events with charging stations. Perhaps a faulty charger is causing over-voltage spikes." This shows systems thinking – looking at not just the car but the ecosystem (charging network, grid). - **Edge Deployment Consideration:** Note that some processing might happen in the vehicle. You could mention, "In a real deployment, some anomaly detection could run on an edge device (in-car) to reduce data upload needs, and only send summary or alerts to the cloud." This demonstrates awareness of bandwidth and latency constraints. For example, it's not feasible to stream 500ms data from *millions* of cars continuously to the cloud due to cost; instead edge computing can pre-filter important data. If you mention technologies like AWS Greengrass or Azure IoT Edge it shows you're thinking ahead. - **Compliance & Security:** Consider adding a brief note in your README about data privacy: e.g. "All sensitive data (like VIN or GPS) can be anonymized or handled according to GDPR guidelines". Or mention role-based access if relevant (maybe only authorized users can see data – though for a demo, you might have a simple login). Security is huge in automotive (they worry about hackers spoofing telemetry or invading privacy). If you can say you're aware of basic security practices (using HTTPS, encrypting data at rest maybe, etc.), that's a nice professional touch.

**WSL2 Cross-Platform Skills in Interviews:** You asked if you should emphasize this – yes, but frame it in terms of Linux competency. The fact that you used WSL2 means you effectively worked in Linux (Ubuntu 22.04) for all development (the code ran in Linux, you used Linux package managers, etc.). That is very valuable because many companies will deploy on Linux servers or use Linux containers. It also shows adaptability – you didn't stick to pure Windows/.NET or something; you ventured into the Unix world, which is where a lot of infrastructure lives. So definitely mention: *"My development was done on Windows via WSL2, giving me a full Linux environment for Docker and tools – so I'm comfortable in both Windows and Linux."* Some employers might ask if you know PowerShell or Bash; you can say you used Bash in WSL2 extensively (for git, docker-compose, etc.). Cross-platform knowledge also means you could help ensure software runs on different OSes (for example, if someone tries to run your project on Mac or Linux, it should work via Docker just as it did on Windows/WSL).

One caveat: don't oversell WSL2 itself as a skill – not all know what it is. Instead, use it to demonstrate *Linux skills and container skills*, as above. If the job environment is primarily Linux, they might even ask if you're comfortable using Linux as your main OS – you can answer that WSL2 experience has basically trained you for that.

**Portfolio Presentation:** When you demo this project (maybe in an interview or via your portfolio site), structure it like a story: 1. Start with a scenario: e.g. "Imagine a fleet of EVs sending data to our system; here we see one vehicle's battery performance over a day..." 2. Show the UI, point out an anomaly, label it, export it. 3. Then perhaps show a snippet of code or architecture diagram to discuss how data flows from ingestion to database to visualization. 4. If asked about decisions (Timescale vs X, etc.), you have all the reasoning we discussed.

Make sure to have the **sources** of your research at hand (you might not need to cite them in the interview like an essay, but the knowledge you've gathered – e.g. Tesla's 500ms, China's GB/T, EU battery passport – can be sprinkled in conversation to show you're well-read). For instance: *"Interestingly, China already requires EVs to upload battery data to a central system – so I built my project in a way that could support such data sharing if needed."* Comments like that can really make you memorable as someone who understands the global context.

**Regulatory Considerations (Future-Proofing):** It's worth summarizing those in conversation: - Mention the **EU Battery Passport** coming in 2026 which will require tracking of battery metadata for recycling/resale <sup>29</sup> <sup>30</sup> . Say how a system like yours could feed into the data needed for a battery passport (like providing the usage history and health metrics that get attached to the passport). - Mention data privacy and the concept of owner consent for sharing vehicle data – the industry is moving toward giving owners access to their own telemetry (there are laws in EU for it). Your system could be seen as a tool for owners or service centers, which implies you'd build in user authentication, etc., down the road. - **Cybersecurity:** Car telemetry systems must be secure to prevent malicious actors injecting false data or breaching vehicles. You can mention you would secure the API (Auth tokens, encryption) and that you're aware of standards like **UNECE WP.29** requiring cybersecurity management for vehicles. This might be beyond what a portfolio project implements, but mentioning it is a differentiator.

**Edge and Federated Trends:** If asked about the future, highlight how you foresee more computation happening on the car (edge AI for battery diagnostics), and data being aggregated in a privacy-preserving way (federated learning – e.g. cars share model updates, not raw data). This shows you're not just thinking of the status quo but the next 5-10 years, which many employers will find impressive in a candidate.

To concisely answer the specific quick questions you posed:

1. **Is PostgreSQL/TimescaleDB choice defensible for 2024-2025?** – Yes, absolutely. It's a modern choice that demonstrates SQL and time-series knowledge. Many companies are adopting similar stacks. You can cite its performance and successful use cases (as we discussed) in defense <sup>19</sup> . No interviewer should ding you for using Timescale; if anything, they'll be interested to hear why and how it went.
2. **Streaming components (Kafka/Redis) vs simplicity?** – For the MVP, it's wise to keep it simple (focus on core functionality and make it solid). Adding Kafka just for show could introduce complexity and risk. Instead, you can simulate streaming and *talk about* how you'd add Kafka or Redis Streams if scaling up. Most interviewers will appreciate a clean design over a buzzword-laden but fragile prototype. However, be ready to discuss event-driven design. If you have extra time, implementing a small WebSocket live update feature (without full Kafka) could be a middle ground to demonstrate real-time capability (e.g. stream the latest data point to the UI every second from FastAPI). But this is sugar on top – the main value is in data correctness and analysis.
3. **What would make it stand out?** – The domain-specific insights and polish. Concretely: the interactive anomaly labeling (few candidates will have built something like that), the visualization quality, and the depth of metrics analyzed (not just generic "temperature vs time", but relating it to EV-specific phenomena). Also, documentation can set you apart – provide a short report or README that includes an architecture diagram, explains choices, and maybe references standards (showing you did research). Many portfolio projects lack thorough documentation; yours can shine here.



Lastly, showing results – e.g. “Here we see this vehicle’s battery is degrading faster, likely due to frequent fast charging” – connecting the data to a narrative demonstrates you can generate actionable insight, not just pretty graphs.

4. **Features/use-cases to demonstrate domain knowledge?** – We touched on many: SoH tracking, anomaly examples that mirror real EV issues (thermal runaway, cell imbalance, charger fault), comparison of driving vs charging behavior, mention of second-life or end-of-life decisions using data, etc. Even the inclusion of things like ambient temperature or the effect of high C-rate charging will show you understand EV batteries, not just data pipelines.
5. **WSL2 + Docker + TSDB gotchas/tips?** – Summarizing: Use Linux filesystem for data for performance <sup>22</sup>, monitor resource usage via `.wslconfig` or Docker resources if needed, avoid `localhost` confusion by using Docker DNS (`db:5432` from FastAPI container instead of trying host networking, unless you do host mode), and keep an eye on time sync (WSL2 time can drift if VM is paused, but usually minor). Also, ensure Docker Desktop is set to use WSL2 backend (which it is by default nowadays). One gotcha: if using VS Code with Remote-WSL and also Docker, make sure your Docker context is correct (it should be the default, connecting to Docker Desktop which is already in WSL). If any performance issues occur, a known tip is disabling Windows Defender scanning of the WSL files (it can slow I/O). You can mention that if relevant. But overall, no show-stoppers in WSL2 – it’s production-grade for dev now.
6. **Emphasize cross-platform (WSL2) in interviews?** – Yes, mention that through WSL2 you’ve effectively been developing and testing in a Linux environment, which means you can easily deploy to Linux servers. It shows you are not confined to Windows and that you’ve used modern dev environments. Many companies use WSL2 internally too, so they’ll nod in understanding. It’s not a primary skill on its own, but it reinforces that you have DevOps savvy. For example, “I used Docker and WSL2 so the dev environment is the same as a Linux server, making deployment trivial.” That’s a great point to make for any full-stack or data engineering role (where often the development vs production environment differences cause issues).

## 6. Future-Proofing and Emerging Trends in EV Battery Monitoring

To ensure your project remains relevant and shows forward-thinking, consider these trends and how you can address or mention them:

- **Edge Computing & On-Board Analytics:** As hinted before, more intelligence is moving on-board. Modern EVs (like Tesla) already do a lot of processing in the car – for instance, they might run algorithms to estimate battery parameters in real time and only send summary data or exceptions to the cloud. The rise of powerful automotive processors (NVIDIA Orin, etc.) means cars could run anomaly detection or even minor ML models locally. You could mention in a “Future Work” section that the system could be adapted to run in a hybrid cloud-edge fashion: critical alerts detected on the edge (for immediate driver notification and to reduce bandwidth), with raw data uploaded in batch for deeper cloud analysis. Edge computing is especially important for fleets that might not have constant connectivity – a truck in a remote area might store data and upload later. Designing your system to handle intermittent connections (buffering data, handling out-of-order timestamps) is a robustness angle you can mention.

- **Federated Learning & Privacy:** Companies like Tesla have massive fleets and they want to learn from the data without transmitting absolutely everything (to respect privacy and save bandwidth). Federated learning is an approach where the training of models happens on the edge and only model weights or gradients are shared to a central server (not raw data). While implementing this is beyond scope, you can show awareness: *"If we were building an ML model for battery RUL, we could employ federated learning – sending a model to vehicles, training on local data, and aggregating improvements – which addresses privacy and uses each car's computing power."* This is bleeding-edge in automotive (some research but not widespread yet), so even knowing the term and concept will impress some interviewers.
- **Integration with Big Data and AI Platforms:** The industry is embracing big data platforms like **Databricks** or **Snowflake** for analyzing vehicle data at scale. In the future, your system might export labeled datasets to such platforms for model training. You can mention something like: "Our labeled anomaly dataset could be fed into a Databricks pipeline to train, say, a classifier that predicts thermal runaway risk. We'd then integrate that model either back into our system or deploy it to vehicles." This shows you understand the ML lifecycle – collecting data, labeling, training, deploying – which is exactly what many employers want (even if your project stops at labeling).
- **Scalability & Cloud-native Options:** Future-proofing also means thinking about how the system would scale to say 100k vehicles. You could say: to scale out, you might use a distributed TSDB or a cloud service. For example, use Timescale Cloud or convert to a clustered solution (Timescale can do multi-node now in enterprise versions). Or use Kafka + Spark streaming for processing instead of feeding everything through a single FastAPI instance. Essentially outline that the architecture is modular: you can scale ingestion horizontally by adding broker consumers, you can scale the DB by partitioning by vehicle or using cloud storage for cold data, etc. Knowing how to scale will set you apart from someone who only ever played with toy data.
- **Battery Tech Evolution:** The monitoring needs might evolve with new battery tech (like solid-state batteries or new chemistries). For instance, solid-state might eliminate some thermal runaway issues but might have other failure modes. Being aware of this isn't crucial for a software role, but if it comes up, you can discuss how monitoring might adapt (e.g. maybe less about overheating, more about cell resistance growth).
- **Regulations & Compliance:** We already discussed EU Battery Passport <sup>29</sup>. You can also mention the **EU Data Act** which will enforce that vehicle data be accessible to users and third parties (breaking OEM monopolies on data) <sup>31</sup>. This could mean your system might one day connect to standardized data streams from any vehicle (with user permission) – an exciting prospect for aftermarket or third-party analytics. Car makers themselves are preparing for this by building platforms to share data (e.g. BMW's CarData, or Mercedes's data platform). If you mention your awareness of these, it shows you're in tune with the business/legal landscape, not just coding.
- **Environmental and Safety Monitoring:** ESG (Environmental, Social, Governance) goals mean companies pay more attention to battery lifecycle (recycling, reuse). A future-proof system might provide data for **carbon footprint analysis** of batteries or **traceability** (like which materials have high degradation, indicating manufacturing issues). It might be tangential, but for instance, if you record each battery's metadata (manufacturer, chemistry), your system could help identify if certain battery batches degrade faster – that's valuable info (recently, GM had to recall Bolts due to battery defects; data monitoring could catch issues early). Mentioning this kind of use case – early fault detection leading to recalls or proactive fixes – will show you grasp the real-world impact.

## Concluding Recommendations:

In summary, your EV Battery Health Monitor project is well-aligned with industry practices and needs. PostgreSQL/TimescaleDB provides a solid, modern foundation for time-series data in 2024/2025, balancing the familiarity of SQL with the performance of a specialized TSDB. The chosen tech stack (FastAPI, React, Docker on WSL2) is both current and in-demand, demonstrating full-stack ability which is attractive to employers in automotive tech who often need engineers that can span software layers. By incorporating industry standards – key metrics like SoC, SoH, temperatures, and aligning with protocols (like referring to real OEM APIs or standards) – you show domain credibility, not just programming skills.

Keep the project scope achievable (you have a lot of great ideas; prioritize the core functionalities that you can deliver cleanly in 4 weeks). A polished, well-documented MVP that you deeply understand is better than an overly complex one that only half works. It sounds like the focus will be on **data correctness, interactive visualization, and insightful labeling**. Those are exactly the things that will catch an interviewer's eye – because they can imagine how that would apply to their needs (be it diagnosing an EV fleet's issues or building a cool consumer-facing battery health report).

Lastly, **present confidence in your choices**. If asked "Why didn't you use X?", explain the trade-offs (e.g., "I chose Timescale over InfluxDB for SQL capabilities and easier integration; Influx could also work and I'm familiar with its concepts, but for a single-stack project Timescale made development faster"). This shows you are pragmatic and can justify technical decisions, a key skill in any senior role.

With these research insights and your solid implementation plan, you're set to not only build a great portfolio project but also to discuss it in a way that highlights both your engineering expertise and your understanding of the EV domain. Good luck – it's an exciting and highly relevant project, and your thoughtful approach to aligning it with industry trends will surely pay off in interviews!

### Sources:

- NREL – *Electric Vehicle Lithium-Ion Battery Life Cycle Management*, on China's GB/T telematics standard <sup>1</sup>
- Tesla – *Fleet Telemetry API Documentation*, on telemetry frequency (500 ms) <sup>2</sup>
- Ampcontrol – "10 Essential Electric Vehicle Tracking Metrics for EV Fleets," on key EV telemetry metrics like SoC, kWh, range, alerts <sup>32 11 13</sup>
- AWS Big Data Blog – "Real-time and batch analytics of electric vehicles on AWS," on data volume and pipeline architecture <sup>33 34 14</sup>
- TigerData (Timescale) – *8 Reasons to Choose TimescaleDB*, on time-series performance (90%+ compression, query speedups) <sup>18 19</sup>
- Synnax Labs – *Performance Compared: Synnax vs Timescale vs Influx*, on Timescale handling tens of millions of inserts per second in benchmarks <sup>20</sup>
- MOBI Initiative – *Battery State of Health standard*, defining SoH as ratio of current to original capacity <sup>7</sup>
- Geotab – *EV Battery Degradation Study* (cited via UCS blog), on real-world battery SOH decline (~2.3%/year) <sup>35</sup> (for context on importance of tracking SOH).

---

<sup>1</sup> Electric Vehicle Lithium-Ion Battery Life Cycle Management

<https://docs.nrel.gov/docs/fy23osti/84520.pdf>

2 5 6 10 Available Data | Tesla Fleet API

<https://developer.tesla.com/docs/fleet-api/fleet-telemetry/available-data>

3 4 11 12 13 32 10 Essential Electric Vehicle Tracking Metrics for EV Fleets

<https://www.ampcontrol.io/post/electric-vehicle-tracking-for-ev-fleets>

7 8 9 Standard for EV battery 'state of health' – MOBI | The New Economy of Movement

<https://dlt.mobi/standard-for-ev-battery-state-of-health/>

14 24 25 33 34 Building a scalable streaming data platform that enables real-time and batch analytics of electric vehicles on AWS | AWS Big Data Blog

<https://aws.amazon.com/blogs/big-data/building-a-scalable-streaming-data-platform-that-enables-real-time-and-batch-analytics-of-electric-vehicles-on-aws/>

15 Does anyone offer a list of vehicle telemetry stats? - Reddit

[https://www.reddit.com/r/electricvehicles/comments/vzrc0s/does\\_anyone\\_offer\\_a\\_list\\_of\\_vehicle\\_telemetry/](https://www.reddit.com/r/electricvehicles/comments/vzrc0s/does_anyone_offer_a_list_of_vehicle_telemetry/)

16 17 18 19 21 8 Reasons to Choose Timescale as Your InfluxDB Alternative | TigerData

<https://www.tigerdata.com/learn/reasons-to-choose-timescale-as-your-influxdb-alternative>

20 Performance Compared: Synnax vs. Timescale vs. Influx | Synnax

<https://docs.synnaxlabs.com/guides/comparison/performance/one-billion-rows>

22 23 How I fixed WSL 2 filesystem performance issues – Rob Pomeroy

<https://pomeroy.me/2023/12/how-i-fixed-wsl-2-filesystem-performance-issues/>

26 Kafka Connect Open Source Development | Geotab

<https://www.geotab.com/blog/kafka-connect/>

27 35 EV Battery Health Insights: Data From 10,000 Cars | Geotab

<https://www.geotab.com/blog/ev-battery-health/>

28 EV\_battery\_charging\_data - Kaggle

<https://www.kaggle.com/datasets/ziya07/ev-battery-charging-data>

29 Are you ready for the EU battery passport? Everything you need to ...

<https://www.4flow.com/blog/are-you-ready-for-the-eu-battery-passport-everything-you-need-to-know>

30 Blog - Battery Passports and the Impact of ESG Regulations

<https://www.accure.net/blogs/battery-passports-and-the-impact-of-esg-regulations>

31 Electric Vehicles: A Data-First Revolution in the Automotive Industry

<https://allindiaev.com/electric-vehicles-a-data-first-revolution-in-the-automotive-industry/>